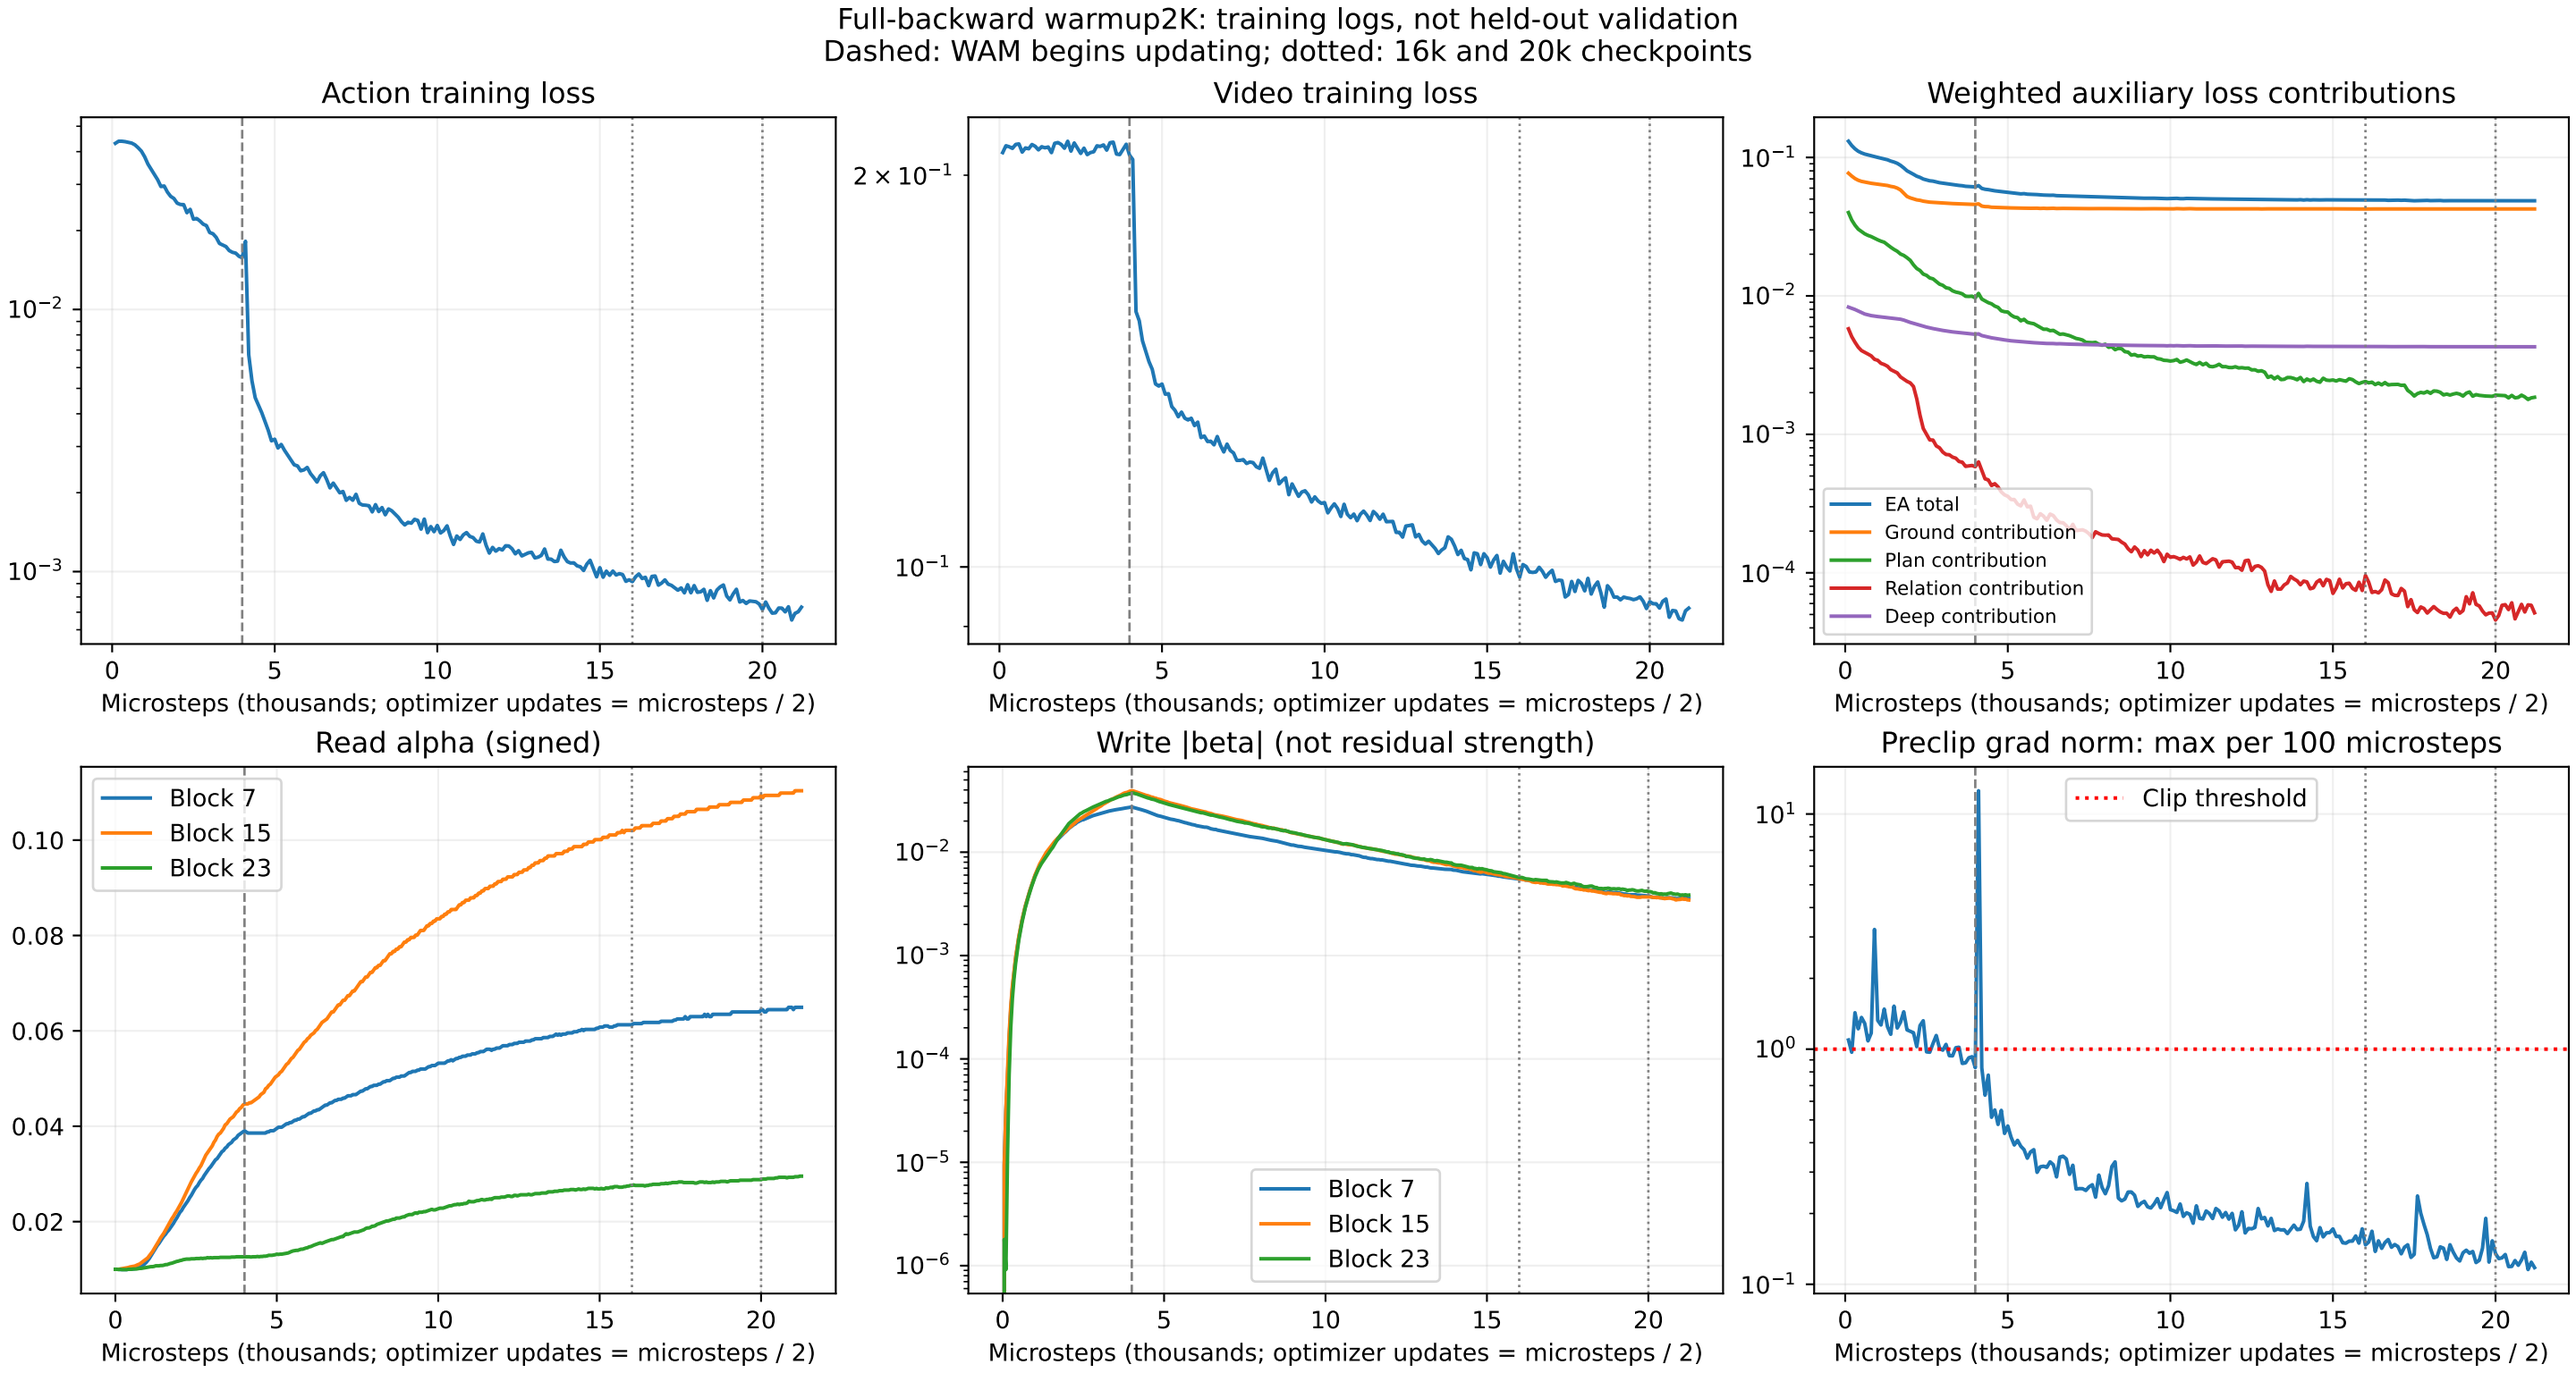
Full-backward warmup2K: training logs, not held-out validation
Dashed: WAM begins updating; dotted: 16k and 20k checkpoints
Action training loss
Video training loss
Weighted auxiliary loss contributions
Read alpha (signed)
Write |beta| (not residual strength)
Preclip grad norm: max per 100 microsteps
Microsteps (thousands; optimizer updates = microsteps / 2)
EA total
Ground contribution
Plan contribution
Relation contribution
Deep contribution
Block 7
Block 15
Block 23
Clip threshold